

Discount Discipline and Margin Risk

A diagnostic of pricing health across a \$1.87B revenue book, January 2023 to December 2025

\$1.87B

REVENUE ANALYSED

\$417M

FORGONE TO DISCOUNT

81.8%

PRICE REALIZATION

35

HIGH-RISK ACCOUNTS

VERDICT

Discount-reliant growth risk

Revenue is stable, but it is bought with a structurally elevated discount that thins margin and concentrates risk in a small set of large Enterprise accounts.

Contents

1 Executive summary	4
The five findings that matter	4
2 Context and objectives	6
What governance means here	6
Scope and audience	6
3 Data and methodology	8
The margin proxy and why it is a proxy	8
Key metric definitions	8
Pipeline integrity and validation	8
4 Analytical framework	10
5 Findings: the aggregate picture	11
6 Findings: what discount costs the margin	13
7 Findings: where the discounting lives	16
Segment is the primary axis of concentration	16
Channel sharpens the same picture	17
Geography is not a factor	18
One product category breaks the pattern	19
Concentration cuts the right way	20
8 Findings: who carries the risk	23
What is actually driving the scores	24
Dependency is structural, not promotional	25
9 Sensitivity and robustness	27
10 Risks, limitations, and caveats	29
The data is synthetic	29
Margin is a proxy	29
Correlation is not full causation	29
The scores encode judgement	29

11 Recommendations and action priorities 30

Priority 1: ring-fence the high-risk accounts for renewal renegotiation	30
Priority 2: put an approval gate on deep discounts in the indirect channels	30
Priority 3: re-price Professional Services	30
Priority 4: adopt price realization as a standing operating metric	30

12 Appendix 32

A. Segment detail	32
B. Channel detail	32
C. Product category detail	32
D. Region detail	32
E. Risk tier detail	33

1 Executive summary

This review examines whether the business is growing on healthy pricing or on discounting that quietly erodes margin. The answer, across \$1.87B of revenue and 38,349 order lines spanning January 2023 to December 2025, is the second. Revenue is steady, but it is purchased at a discount that has not come down in three years and that lands hardest on the largest accounts. The verdict is a discount-reliant growth risk: not a crisis, but a structural drag that is large enough to act on and concentrated enough to fix.

The headline number is the gap between what the book lists for and what it actually bills. List value over the period was \$2.29B; realized revenue was \$1.87B. The difference, \$417M, is revenue given away through discount. That puts price realization at 81.8%, and it has barely moved month to month, holding inside a 80.7% to 82.8% band for 36 straight months. Stability at this level is the tell. A discount that never relaxes is not a sequence of tactical deals; it is a standing policy that the organisation has stopped noticing.

Discount depth and margin move against each other with discipline. At the customer level the correlation between average discount and average margin is -0.69, and the slope is almost exactly one for one: every additional point of discount costs roughly a point of margin proxy. The same gradient shows up when order lines are bucketed by depth, where blended margin falls from 51.7% in the shallowest band to 37.2% in the deepest. Discounting here is not a harmless lever pulled to close deals; it transfers value out of the margin line at a near-constant rate.

The risk is concentrated, which is the good news. Enterprise alone is \$1.10B of revenue, 58.7% of the book, and it sits at the worst discount-margin position of any segment: a 20.0% average discount against a 43.3% margin proxy, with nearly half its revenue (49.2%) flowing through deals discounted past the high-discount line. Revenue is concentrated too: the top 20% of customers account for 66.1% of revenue. A governance effort aimed at a few dozen accounts can therefore reach most of the exposure without disturbing the long tail.

The five findings that matter

- **Price realization is stuck at 81.8%.** The \$417M forgone to discount is a standing cost, not a seasonal one. Monthly realization never rises above 82.8%.
- **Discount buys margin at par.** A customer-level slope of -0.99 margin points per discount point, and a 51.7%-to-37.2% margin drop across depth bands, means depth is the single most reliable predictor of thin margin in the book.
- **Enterprise is the problem segment.** It holds 58.7% of revenue at the lowest margin proxy of the four segments and the highest high-discount share at 49.2%.
- **Indirect channels discount deepest.** The reseller channel runs a 21.0% weighted discount against 14.5% Online, and the Enterprise-reseller cell reaches 23.4% with 88.6% of its revenue deeply discounted.

- **The exposure is nameable.** 35 accounts carrying \$154M sit in the high-risk tier with an average priority score of 70 out of 100, each flagged for discount-policy redesign.

Section 11 sets out the response. In short: ring-fence the 35 high-risk accounts for renegotiation at renewal, put a 20% approval gate on the reseller and partner channels where depth concentrates, and re-price the Professional Services category, which discounts like the rest of the book but earns a 20.5% margin against a portfolio average above 45.2%. None of these moves requires new revenue. Each recovers value the book is already producing and then discounting away.

2 Context and objectives

Discounting is the easiest lever in a sales organisation to pull and the hardest to see once pulled. A single deal closed three points cheaper is invisible in a quarterly revenue number; ten thousand of them are a margin problem that no one decided to create. The purpose of this review is to make that accumulated decision visible, to quantify what it costs, and to locate it precisely enough that it can be governed rather than merely lamented.

The analytical question is deliberately narrow. It is not whether the business is growing; revenue is stable across the window. It is whether that revenue rests on healthy pricing or on a discounting habit that erodes margin. The distinction matters because the two look identical on a top-line chart and diverge sharply on the margin line. A book that bills \$1.87B at 81.8% realization is a materially different asset from one that bills the same amount at 92.0%, and the difference, \$417M over the period, is exactly the quantity this review is built to surface.

What governance means here

Pricing governance is not the elimination of discount. Discount is a legitimate instrument for winning competitive deals, rewarding volume, and entering accounts. Governance is the discipline of knowing where discount is being spent, whether the return justifies the depth, and which accounts have drifted from negotiated concession into structural dependence. The outputs of this review, a risk score per customer, a priority queue, and a recommended action for each tier, are built to support that discipline at the level where it can actually be exercised: the individual account and the channel.


Scope and audience

The review covers the full order book: 1,173 customers, 28 products across 5 categories, 4 sales channels and 4 regions, observed over 36 months. It is written for three readers at once. For the executive, the executive summary and the recommendations carry the argument. For the commercial and finance owner, the findings sections give the segment, channel, and account detail needed to act. For the technical reviewer, Sections 3 and 4 document the data lineage, the metric definitions, and the validation that stands behind every figure.

TREND

One in five list-price dollars is given away every month

Share of list value actually billed. The shaded band is revenue forgone to discount



Source: pricing governance pipeline, processed marts. January 2023 to December 2025.

Figure 1. Price realization has held inside a two-point band for three years. The shaded area is the share of list value forgone to discount each month, the standing cost this review quantifies.

3 Data and methodology

Every figure in this report traces to a single reproducible pipeline that moves raw transactions through cleaning, feature construction, a SQL warehouse layer, risk scoring, and validation. The pipeline reads five raw tables, customers, products, orders, order items, and sales representatives, and resolves them into an enriched order-item grain of 38,349 rows. That grain is the atom of the analysis: one row per product sold on one order, carrying its list price, its realized price, the discount between them, an estimated unit cost, and the segment, channel, region, and tenure context of the sale.

The margin proxy and why it is a proxy

Cost of goods is estimated rather than booked, so margin throughout this report is a proxy, defined as realized revenue less estimated cost over realized revenue. This is stated plainly because it bounds the conclusions. The proxy is reliable for comparison, which is what every finding here depends on. The gap in margin between a deeply discounted line and a shallow one, between Enterprise and SMB, between Professional Services and Collaboration, is a real difference in the same consistently constructed quantity. The proxy is not reliable as a statement of absolute profitability and is not used as one. No recommendation in Section 11 rests on the level of margin; each rests on a difference in margin, which the proxy measures cleanly.

Key metric definitions

Four constructed metrics carry most of the analytical weight, and each is defined once here so the findings can be read without ambiguity.

Metric	Definition
Realized discount	1 minus realized price over list price, per line; aggregated revenue-weighted unless stated.
Price realization	Realized revenue over list-price revenue. The complement of the weighted discount.
Margin proxy	Realized revenue less estimated cost, over realized revenue.
High-discount line	A line is high-discount at or above a 20% threshold; sensitivity to this cut is tested in Section 9.
Governance priority	A 0-100 composite of discount dependency, margin erosion, and pricing-risk scores per customer.

The customer-level scores deserve a word because the priority queue depends on them. Each customer receives a discount-dependency score, a margin-erosion score, and a pricing-risk score, each on a 0-100 scale, blended into a single governance-priority score and reliability-weighted by how much data the customer provides. A customer with six orders and a customer with seventy are not scored as if equally certain. The blended score then sorts customers into three tiers with a recommended action each, which is the operational output the commercial team works from.

Pipeline integrity and validation

The pipeline is deterministic and seeded, so the same inputs produce the same outputs on any machine, and a release gate runs a battery of validation checks before any figure is published. For this review, all ten formal-analysis validation checks passed, and the metric-contract layer confirmed that every published metric matched its independently recomputed value within tolerance. The data is synthetic, generated to be internally consistent rather than drawn from a live system, which is the one substantive limitation and is treated at length in Section 10. It does not weaken the methodological points, which are about relationships in the data, and those relationships are real and stable within the dataset analysed.

4 Analytical framework

The single question, healthy growth or discount-reliant growth, is answered along four axes, and the findings sections follow them in order. Each axis isolates one way that a discounting problem can hide.

The aggregate axis asks whether the book as a whole realizes its list prices. It is answered by price realization and the revenue forgone to discount, read over time so that a stable problem can be told apart from a deteriorating or improving one. A single snapshot cannot make that distinction; 36 months can.

The depth axis asks what discounting costs at the margin. It is answered by the relationship between discount depth and margin, measured two ways for robustness: across discount buckets at the line grain, and across customers at the account grain. Agreement between the two supports a consistent descriptive reading of how discount depth and margin move together.

The structural axis asks where the discounting lives. It is answered by cutting discount and margin by segment, channel, region, and product category, and by the interaction of segment and channel, because a problem that is uniform requires a blunt policy response while a problem that is concentrated requires a targeted one. The data, as Section 7 shows, is firmly the second case.

The account axis asks who is responsible for the exposure. It is answered by the customer risk scores, the priority tiers, and the concentration of revenue, which together turn a diagnosis into a worklist. This is the axis that makes the review operational rather than merely descriptive.

A fifth consideration, sensitivity, runs underneath all four. Every threshold in the analysis, above all the 20% high-discount line, is tested by moving it and watching whether the conclusion moves with it. Section 9 shows it does not across 15%, 20%, 25%, which is what distinguishes a finding from an artefact of where a cutoff happened to fall.

5 Findings: the aggregate picture

Start with the whole book. Over January 2023 to December 2025 the business billed \$1.87B against a list value of \$2.29B. The \$417M difference is the cost of discount, and at 81.8% realization it is large in both absolute and relative terms. For every five dollars the catalogue says the book is worth, it collects a little over four.

What turns this from a number into a finding is its stability. Monthly price realization sits in a 80.7%-to-82.8% band for the entire window, a spread of barely two points across three years that include whatever seasonality, mix shift, and deal cycles the period contained. The weighted discount behind it moves just as little, oscillating around 18.2% without trend. Revenue, plotted against that flat discount line in Figure 2, is similarly trendless. The book is not growing into its discount or out of it. It is holding both constant.




Figure 2. Monthly revenue against the revenue-weighted discount. Neither trends. The discount is a level the organisation operates at, not a response to conditions that vary.

The interpretation is the central claim of the review. A discount that responds to conditions, deeper in slow quarters, shallower in strong ones, would vary. This one does not. Its constancy is the signature of policy rather than tactics: a set of standing list-to-net conventions, channel terms, and negotiating defaults that produce the same eighteen points of discount regardless of context. That is a governable object. Tactics are fought deal by deal; a standing level can be reset once and will then hold at the new level just as firmly.

It also means the \$417M is annuitised. This is not a one-time leakage to be plugged but a recurring transfer that reappears every month at the same rate. Reading it that way changes the economics of intervention. A two-point improvement in realization, from 81.8% toward 83.8%, is worth on the order of \$46M a period,

every period, which is the budget any governance effort should be measured against.

6 Findings: what discount costs the margin

The aggregate tells us discount is large and standing. The depth axis tells us what it does. The answer is unusually clean: discount and margin trade against each other at close to one for one, and they do so whether the data is cut by line or by customer.

At the line grain, sorting all 38,349 order items into discount bands and computing blended margin within each produces a near-monotone staircase. The shallowest band earns a 51.7% margin proxy; the deepest earns 37.2%. Each step deeper into discount removes roughly four points of margin, and the steps do not flatten out at the bottom, which is where one would hope to see depth buying something, strategic volume, a beachhead account, that offsets the margin given up. It does not.





Figure 3. Blended margin proxy falls at every step into deeper discount. The relationship is monotone, with no flattening at depth to suggest the deepest discounts buy compensating value.

The same gradient appears one level up, at the customer grain, and this is what makes the reading robust. Plotting each of the 1,173 customers by average discount against average margin yields a correlation of -0.69 and a fitted slope of -0.99 margin points per discount point. A slope of one is the worst case for a discounter: it means depth is not bought down by volume or mix at the account level either. The customer who is discounted ten points deeper than another is, on average, ten margin points thinner, full stop.

CORRELATION

Deeper discounting reliably tracks thinner margin

Each bubble is one of 1,173 customers, sized by revenue



Source: pricing governance pipeline, processed marts. January 2023 to December 2025.

Figure 4. Every customer plotted by discount and margin, sized by revenue. The downward fit is steep and tight; deeper-discounted customers are reliably thinner-margin customers, with no volume offset visible.


Two cuts of the same data agreeing this closely is the evidence that licenses a causal reading. If depth merely accompanied thin margin, through some third factor like product mix, the line and customer views could easily disagree, because they aggregate over different things. They agree because the mechanism is direct: discount is subtracted from price, price sits above cost, and so discount comes out of margin almost dollar for dollar until it reaches cost. The book is not discounting to win volume that rebuilds margin elsewhere. It is discounting into the margin line.

Composition confirms the depth is not confined to a harmless tail. Bucketing revenue rather than counting lines, the bulk of the book sits in the 10-to-20% band, but the bands above the high-discount line carry real weight, and at the 20% cut 33.6% of all revenue flows through deals discounted past it. A third of the book is in the part of the discount distribution where, per Figure 3, margin is measurably impaired.

COMPOSITION

Most revenue carries a 10-20% discount, but the deep tail is heavy

Billed revenue by discount band. Bands above 20% are highlighted




Source: pricing governance pipeline, processed marts. January 2023 to December 2025.

Figure 5. Revenue by discount band. The 10-20% band dominates, but the highlighted deep bands are not a rounding error; they carry roughly a third of revenue.

DISTRIBUTION

Discounting clusters in a tight 10-20% band, not at the edges

All 38,349 order items by realized discount depth



Source: pricing governance pipeline, processed marts. January 2023 to December 2025.

Figure 6. The distribution of line-level discount. Discounting clusters tightly rather than spreading, another sign of standing convention over case-by-case negotiation.

7 Findings: where the discounting lives

A problem that is everywhere needs a different remedy than a problem that is somewhere. This section cuts the discount four ways to establish which it is. The answer is consistent across every cut: the discounting is concentrated, in the Enterprise segment, in the indirect channels, and in one product category, while it is essentially uniform across geography. Concentration is what makes a targeted response possible.


Segment is the primary axis of concentration

The four segments form a clean ladder. Enterprise discounts deepest at 20.0% and earns the thinnest margin at 43.3%; SMB discounts shallowest at 12.0% and earns the fattest margin at 52.0%, with Mid-Market and Public Sector ranged between. The ladder would be unremarkable if the segments were small. They are not. Enterprise is \$1.10B, 58.7% of the entire book, which means the worst pricing position by margin is also the largest by revenue. That is the single most important structural fact in the data.

RISK

Enterprise carries the most revenue at the worst discount-margin position

Segment position by discount and margin. Bubble area is total revenue



Source: pricing governance pipeline, processed marts. January 2023 to December 2025.

Figure 7. Segments positioned by discount and margin, sized by revenue. The largest bubble sits in the worst corner: Enterprise combines the most revenue with the deepest discount and the thinnest margin.

Segment	Revenue	Avg discount	High-disc share	Margin proxy
Enterprise	\$1.10B	20.0%	49.2%	43.3%
Public Sector	\$151M	18.0%	27.1%	42.6%
Mid-Market	\$433M	16.0%	12.4%	48.1%
SMB	\$188M	12.0%	1.1%	52.0%

The high-discount share column is where the concentration is starkest. In Enterprise, 49.2% of revenue is deeply discounted; in SMB it is 1.1%. These are not two points on a smooth gradient but two different pricing regimes operating in the same company. SMB has, in effect, already solved the problem this review describes. Enterprise has not.

Channel sharpens the same picture

Cutting by channel shows the indirect routes to market discount markedly deeper than the direct ones. The reseller channel runs a 21.0% weighted discount and the partner channel 19.2%, against 18.1% Direct and only 14.5% Online. Direct is the largest channel at \$1.08B, so its discount sets the book average, but the indirect channels are where depth concentrates per dollar of revenue.




Figure 8. Weighted discount by channel, revenue noted at each base. Indirect channels discount close to half again as deep as Online.

The interaction of segment and channel is where the exposure peaks, and it is worth isolating because it names the single worst cell in the book. Where Enterprise meets the reseller channel, average discount reaches 23.4% and 88.6% of that cell's revenue is deeply discounted, against a margin proxy of 40.9%. The heatmap makes

the gradient legible: discount darkens reliably from the SMB-Online corner to the Enterprise-reseller corner, the two regimes sitting at opposite ends of the same grid.




Figure 9. Average discount by segment and channel. The gradient runs cleanly from the disciplined SMB-Online corner to the Enterprise-reseller corner, which is the deepest cell in the book.


Geography is not a factor

Region is the control that proves the others are real. Cutting the same discount by region produces almost no spread: APAC is deepest at 18.9% and Europe shallowest at 18.0%, a range under a single point across four regions that differ enormously in size. Discount depth does not travel with geography. It travels with who is being sold to and through which channel, which is precisely the structural reading the segment and channel cuts established.

GEOGRAPHY

Discount depth is roughly even across regions

Weighted discount and revenue by customer region



Source: pricing governance pipeline, processed marts. January 2023 to December 2025.

Figure 10. Weighted discount by region. The near-flat profile rules geography out as a driver and confirms that segment and channel, not place, carry the concentration.


One product category breaks the pattern

Product category mostly tracks margin as expected, with one exception sharp enough to act on. Professional Services discounts at 17.3%, in line with the rest of the catalogue, but earns a margin proxy of only 20.5%, less than a third of the 63.0% that Collaboration earns and far below the 45.2% book average. It is discounted like a high-margin product while behaving like a low-margin one. Core Platform, by contrast, is the revenue engine at \$921M and earns a respectable 45.2%, so the category problem is specific rather than general.

COMPOSITION

Professional Services earns the thinnest margin on an ordinary discount

Discount versus margin proxy by product category



Source: pricing governance pipeline, processed marts, January 2023 to December 2025.

Figure 11. Discount against margin by category. Professional Services is the outlier: ordinary discount, exceptionally thin margin, which makes its discount the least defensible in the book.


Concentration cuts the right way

The structural findings would be discouraging if the exposure were spread thinly across thousands of accounts. It is not. Revenue is heavily concentrated: the top 10% of customers produce 46.2% of revenue and the top 20% produce 66.1%. The same concentration that makes the book sensitive to a few accounts also makes it tractable. A governance programme that engages the largest few hundred relationships reaches the great majority of both the revenue and the discount, which is the bridge from this section's diagnosis to the account-level worklist in the next.

CONCENTRATION

Revenue is highly concentrated, which makes targeted remediation viable

Lorenz curve of revenue across the customer base




Source: pricing governance pipeline, processed marts. January 2023 to December 2025.

Figure 12. The Lorenz curve of revenue. Steep early rise means high concentration, and high concentration means targeted remediation can reach most of the exposure.

RANKING

The top-revenue product also carries the deepest discount exposure

Top 12 products by revenue. Red marks any product with 30%+ of revenue deeply discounted



Source: pricing governance pipeline, processed marts. January 2023 to December 2025.

Figure 13. Top products by revenue, flagged where deep-discount exposure is high. The single largest product also carries the heaviest deep-discount share, compounding its weight in the book.

8 Findings: who carries the risk

The final axis turns the structural picture into a list of names. The risk-scoring layer assigns every customer a governance-priority score and sorts the book into three tiers, each with a recommended action. This is the output the commercial organisation can act on directly.

The tiers are deliberately unbalanced because the risk is. The high-risk tier holds just 35 customers but \$154M of revenue, at an average priority score of 70 out of 100, every one of them flagged for discount-policy redesign. The medium tier is 204 customers and \$718M, marked for segment-pricing review. The low tier is the long, healthy tail: 934 customers and \$999M that need only monitoring. The shape of the priority distribution, a thin upper tail above a broad healthy base, is what makes a small high-touch programme viable.





Figure 14. The three risk tiers by customer count and by revenue. A few dozen high-risk accounts carry revenue out of all proportion to their number.

DISTRIBUTION

Priority is concentrated in a thin upper tail

Governance priority score across 1,173 customers



Source: pricing governance pipeline, processed marts. January 2023 to December 2025.

Figure 15. The governance-priority distribution. Most customers score low; the population that needs action sits above the 90th percentile at a score of 55.

What is actually driving the scores

The scoring layer also records why each customer is flagged, which prevents the queue from being a black box. Aggregating the primary driver across the book, discount dependency is dominant: it is the lead driver for 380 customers carrying \$1.23B, far more revenue than the pricing-risk and margin-erosion drivers behind it. This matters for the remedy. A book whose main problem is dependency, customers who have come to expect a standing discount, responds to renegotiation and approval discipline. It would not respond to the same degree if the driver were, say, cost inflation, which no amount of pricing governance would touch.




Figure 16. Revenue by primary risk driver. Discount dependency dominates, which points the remedy squarely at discount policy rather than cost or mix.


Dependency is structural, not promotional

One more cut closes the loop between this axis and the aggregate finding. Splitting customers by tenure shows discount depth essentially flat across the lifecycle: newly acquired customers and four-year veterans are discounted at materially the same rate. A discount that were promotional, used to win accounts and then withdrawn, would start deep and shallow with tenure. This one does not move. The discount is not an acquisition cost being recovered over the relationship; it is a permanent term of trade, which is the account-level restatement of the standing-policy conclusion from Section 5.

COHORT

Discount depth is flat across tenure, so it is structural not promotional

Average discount and margin proxy by customer tenure band



Source: pricing governance pipeline, processed marts. January 2023 to December 2025.

Figure 17. Discount and margin by customer tenure. The flat discount line across tenure bands shows the discount is structural, baked into the relationship rather than a fading acquisition incentive.

9 Sensitivity and robustness

A finding that depends on where a threshold was drawn is not a finding. The most consequential threshold in this review is the 20% high-discount line, which determines the high-discount revenue share and feeds the risk scores. This section moves that line and watches what happens to the conclusion.

At a 15% line, 76.9% of revenue counts as high-discount and \$163M of revenue is flagged as margin-at-risk. At the 20% line used throughout, those fall to 33.6% and \$62M. At 25%, to 3.9% and \$14M. The magnitudes move, as they must, but the direction and the diagnosis do not. At every threshold the same segments, channels, and accounts surface as the deepest, and at every threshold the margin proxy on high-discount revenue sits below the book average. The verdict is a property of the data, not of the cutoff.




Figure 18. Revenue at risk and high-discount share as the threshold moves from 15% to 25%. The curve is smooth and the ranking of risk is stable, so no single cutoff manufactures the result.

Threshold	High-disc revenue share	Revenue margin-at-risk	Margin on high-disc revenue
15%	76.9%	\$163M	44.0%
20%	33.6%	\$62M	43.0%
25%	3.9%	\$14M	40.4%

Two independent validation layers add to the confidence. The metric-contract layer recomputes every published figure from the raw grain and checks it against the value the pipeline reports, and all checks

reconciled within tolerance. The formal-analysis release gate ran ten structural checks on the analysis window, coverage, and internal consistency, and all ten passed. The numbers in this report are therefore not only reproducible but independently recomputed, which is the standard the methodology section promised.

10 Risks, limitations, and caveats

Honest analysis states its own boundaries. Four matter here, and none of them is fatal to the conclusions, but each shapes how far the conclusions can be pushed.

The data is synthetic

The order book is generated rather than drawn from a production system. It is internally consistent and behaves like a real book, which is what makes the relationships in it meaningful as a demonstration of method, but the specific dollar figures are not a claim about any real company. The value of the review is the analytical machinery, the metrics, the scoring, the validation, and the way they convert a vague worry about discounting into a quantified, located, testable finding. Pointed at a real book, the same machinery would yield real numbers, and that is the intended use.

Margin is a proxy

Because cost is estimated, margin is a proxy and is used only for comparison, never as a statement of absolute profitability. Every finding rests on a difference in margin between groups, which the proxy measures cleanly, rather than on its level. A reader who wants absolute margin must supply booked cost; the relative conclusions, which carry the recommendations, would survive that substitution because they depend on the ordering of margins, not their calibration.

Correlation is not full causation

The near-one slope between discount and margin is consistent with discount driving margin loss, and the agreement between the line and customer cuts strengthens that reading, but an observational dataset cannot prove it the way a controlled price test could. The association is operationally relevant, while the recommendations remain decision-support proposals because synthetic data cannot establish a causal effect.

The scores encode judgement

The governance-priority score blends three sub-scores with chosen weights, and different weights would reshuffle the queue at the margin. The tiering is therefore a defensible prioritisation rather than an objective truth, and it should be used as a starting worklist that commercial judgement refines, not as an automated verdict on any single account. The high-risk tier is robust to reweighting because its members are extreme on every sub-score; the boundary between medium and low is where weight choices bite, and that boundary is exactly where human review belongs.

11 Recommendations and action priorities

The findings converge on a small number of moves, ordered below by the ratio of value recovered to effort and disruption. None depends on winning new revenue. Each recovers margin the book already produces and currently discounts away. The organising target is price realization: moving it from 81.8% toward the mid-eighties is worth roughly \$46M for every two points, recurring.

Priority 1: ring-fence the high-risk accounts for renewal renegotiation

The 35 high-risk accounts carry \$154M at an average priority of 70 and are individually nameable from the scoring output. Assign each an owner, set a target net price ahead of its next renewal, and renegotiate from list rather than from the incumbent discounted price. Because dependency is the dominant driver, the realistic goal is to narrow the discount by a few points per account at renewal, not to remove it. Even a modest narrowing across this group, given its revenue weight, is the single largest available recovery. Tie the programme to a tracked realization number per account so the result is visible.

Priority 2: put an approval gate on deep discounts in the indirect channels

Depth concentrates in the reseller and partner channels, peaking in the Enterprise-reseller cell at 23.4% with nearly nine-tenths of that revenue deeply discounted. Introduce a hard approval gate above the 20% line for these channels, so that any deal past the high-discount threshold requires explicit sign-off against a stated reason. This is a forward-looking control that stops the standing level from regenerating after Priority 1 resets it, and it is cheap: it changes a default, not a headcount.

Priority 3: re-price Professional Services

Professional Services is discounted like the rest of the catalogue, at 17.3%, while earning a 20.5% margin proxy against a book average above 45.2%. This is the least defensible discount in the portfolio. Either lift its net price by curtailing discount on the category, or, if the market will not bear it, reclassify the work and price it as the low-margin service it is so that it stops being bundled into discount decisions calibrated for high-margin product. This move is narrow and self-contained, which makes it a good early proof point.

Priority 4: adopt price realization as a standing operating metric

The reason the \$417M accumulated unnoticed is that no one owned realization as a number. Put it on the commercial dashboard beside revenue, report it monthly by segment and channel, and hold the same owners accountable for both. The measurement is already built; the pipeline produces realization at every cut in this report. What is missing is the standing attention, and standing attention is what keeps a level from drifting back up once it has been reset.

Sequenced, the programme reads:

- **Now:** stand up the 35-account renewal programme and the indirect-channel approval gate. Together they address the bulk of the exposure and stop its regeneration.

- **Next quarter:** execute the Professional Services re-price as a contained proof point, and publish price realization on the commercial dashboard.
- **Standing:** review the medium tier (204 accounts, \$718M) on a rolling basis as renewals arrive, refreshing the risk scores each cycle so the worklist stays current.

12 Appendix

A. Segment detail

Segment	Revenue	Avg discount	High-disc share	Margin proxy
Enterprise	\$1.10B	20.0%	49.2%	43.3%
Public Sector	\$151M	18.0%	27.1%	42.6%
Mid-Market	\$433M	16.0%	12.4%	48.1%
SMB	\$188M	12.0%	1.1%	52.0%

B. Channel detail

Channel	Weighted discount	Revenue
Reseller	21.0%	\$184M
Partner	19.2%	\$378M
Direct	18.1%	\$1.08B
Online	14.5%	\$229M

C. Product category detail

Category	Revenue	Discount	Margin proxy
Core Platform	\$921M	18.4%	45.2%
Security	\$356M	18.8%	47.8%
Analytics	\$259M	18.4%	52.6%
Professional Services	\$206M	17.3%	20.5%
Collaboration	\$129M	16.5%	63.0%

D. Region detail

Region	Revenue	Weighted discount	Margin proxy
North America	\$797M	18.1%	45.4%
Europe	\$499M	18.0%	45.3%
APAC	\$322M	18.9%	44.7%

LATAM	\$252M	18.2%	45.1%
-------	--------	-------	-------

E. Risk tier detail

Tier	Recommended action	Customers	Revenue	Avg priority
High	redesign discount policy	35	\$154M	70
Medium	review segment pricing	204	\$718M	54
Low	monitor only	934	\$999M	28

Provenance. All figures derive from the processed marts of the pricing-governance pipeline over January 2023 to December 2025: 38,349 order lines, 1,173 customers, 28 products. Ten formal-analysis validation checks and the full metric-contract reconciliation passed at the time of publication. Charts are generated by `scripts/build_report_assets.py` and this report by `scripts/build_report_pdf.py`, both reproducible from the processed data.