

Discount discipline and margin risk

A diagnostic of pricing health across a \$1.88B revenue book

Miguel Fidalgo Martins

January 2023 to December 2025

\$1.88B
REVENUE ANALYSED

\$419M
FORGONE TO DISCOUNT

81.8%
PRICE REALIZATION


66
HIGH-RISK ACCOUNTS

Verdict

Realization control gap

Revenue is stable, but list-to-net discipline is leaking value into Enterprise accounts and indirect-channel deals.

Methodology in Section 3 · Limitations in Section 10
All figures reconciled against governed pipeline outputs



Contents

1 Executive summary	4
The five findings that matter	4
Management response at a glance	5
2 Context and objectives	6
What governance means here	6
Scope and audience	6
Decision frame	6
3 Data and methodology	8
The margin proxy and why it is a proxy	8
Key metric definitions	8
Pipeline integrity and validation	9
4 Analytical framework	10
5 Findings: the aggregate picture	11
6 Findings: what discount costs the margin	13
7 Findings: where the discounting lives	16
Segment is the primary axis of concentration	16
Channel sharpens the same picture	17
Geography is not the primary discriminator	19
One product category breaks the pattern	20
Concentration cuts the right way	21
8 Findings: who carries the risk	24
What is actually driving the scores	25

Discount depth persists across tenure cohorts	26
9 Sensitivity and confidence	28
10 Risks, limitations, and caveats	30
The data is synthetic	30
Margin is a proxy	30
Correlation is not full causation	30
The scores encode judgement	31
11 Recommendations and action priorities	32
Priority 1: ring-fence the high-risk accounts for renewal renegotiation	32
Priority 2: put an approval gate on deep discounts in the indirect channels	32
Priority 3: re-price Professional Services	32
Priority 4: adopt price realization as a standing operating metric	32
Control model	32
90-day mobilisation plan	33
Further questions before scale	33
12 Appendix	34
A. Segment detail	34
B. Channel detail	34
C. Product category detail	34
D. Region detail	35
E. Risk tier detail	35

1

Executive summary

The commercial book is not failing; it is under-realising. Across \$1.88B of revenue and 38,391 order lines over January 2023 to December 2025, the business converts only 81.8% of list value into billed revenue. That gap is not random deal noise. It is stable over time, concentrated in Enterprise and indirect channels, and large enough to manage as a value-recovery programme rather than a broad price increase.

The value bridge is clear. List value over the period was \$2.30B; realized revenue was \$1.88B. The difference, \$419M, is the list-to-net concession embedded in the book. Monthly realization stays inside a narrow 80.9% to 82.5% band for 36 months. A metric that stable is not a series of isolated commercial exceptions. It is a standing operating level.

The margin signal is equally direct. At customer level, average discount and average margin carry a -0.68 correlation, with margin falling by an estimated 0.96 points for each additional discount point. The line-level view points the same way: blended margin falls from 51.7% in the shallowest discount band to 38.0% in the deepest. The evidence supports a practical conclusion: depth is the strongest visible marker of margin pressure in this book.

The concentration makes the problem actionable. Enterprise is \$1.10B of revenue, or 58.6% of the book, and carries the weakest discount-margin position: 20.1% average discount, 43.4% margin proxy, and 49.6% of revenue past the high-discount line. The top 20% of customers account for 66.3% of revenue, so a focused programme can reach the value pool without placing the whole sales motion under new controls.

The management implication is specific. Do not treat this as a blanket price-rise mandate. The value pool sits in three places: Enterprise renewal economics, indirect-channel approval discipline, and Professional Services pricing. A two-point realization lift is worth about \$46M on the analysed list-price base, equivalent to roughly \$15M per year at the current run rate, before any volume response. That is the prize to test, not assume.

The five findings that matter

- **Price realization is stuck at 81.8%.** The \$419M forgone to discount is a standing cost, not a seasonal one. Monthly realization never rises above 82.5%.
- **Discount depth marks margin pressure.** A customer-level slope of 0.96 margin points per discount point, and a 51.7%-to-38.0% margin drop across depth bands, makes depth the clearest pricing-risk signal in the book.
- **Enterprise is the priority segment.** It holds 58.6% of revenue at the lowest margin proxy of the four segments and the highest high-discount share at 49.6%.
- **Indirect channels discount deepest.** The reseller channel runs a 21.0% weighted discount against 14.4% Online, and the Enterprise-reseller cell reaches 23.4% with 87.9% of its revenue deeply discounted.
- **The exposure is nameable.** The high-risk tier holds 66 accounts carrying \$267M, at an average priority score of 71 out of 100, each flagged for discount-term review.

Management response at a glance

Move	Scope	Success metric
Renewal reset	66 high-risk accounts carrying \$267M	Account realization lift versus pre-set renewal target
Indirect-channel gate	Reseller and Partner deals above the 20% line	High-discount revenue share and approval exception rate
Services re-price	Professional Services, \$205M revenue at 20.6% margin proxy	Category realization and delivery-margin recovery
Operating cadence	Monthly realization by segment, channel, category and tier	Move book realization toward 83.8%

Section 11 turns those moves into a 90-day programme: ring-fence the 66 high-risk accounts for renewal renegotiation, put a 20% approval gate on reseller and partner deals where depth concentrates, and re-price Professional Services, which earns a 20.6% margin against a portfolio average above 45.3%. None of these actions depends on new revenue. Each is aimed at retaining more of the value the book already creates.

2

Context and objectives

Discounting is easy to approve deal by deal and hard to govern once it becomes the default. One concession is commercial judgement; thousands of similar concessions become a pricing system. This review makes that system visible, quantifies the value at stake, and locates the exposure at the level where management can act: segment, channel, product category and account.

The analytical question is narrow by design. It is not whether the business is growing; revenue is stable across the window. It is whether that revenue is being earned with healthy list-to-net discipline. A book that bills \$1.88B at 81.8% realization is a different asset from one that bills the same amount at 92.0%. The difference, \$419M over the period, is the economic gap this review sizes and prioritises.

What governance means here

Pricing governance is not the elimination of discount. Discount is a legitimate instrument for winning competitive deals, rewarding volume, and entering accounts. Governance is the discipline of knowing where discount is being spent, whether the return justifies the depth, and which accounts have drifted from negotiated concession into structural dependence. The outputs of this review, a risk score per customer, a priority queue, and a recommended action for each tier, are built to support that discipline at the level where it can actually be exercised: the individual account and the channel.

Scope and audience

The review covers the full order book: 1,173 customers, 28 products across 5 categories, 4 sales channels and 4 regions, observed over 36 months. It is written for three readers at once. For the executive, the executive summary and the recommendations carry the argument. For the commercial and finance owner, the findings sections give the segment, channel, and account detail needed to act. For the technical reviewer, Sections 3 and 4 document the data lineage, the metric definitions, and the validation that stands behind every figure.

Decision frame


The report is designed to support three management decisions: where to intervene first, which discounting behaviours should move from field discretion into policy, and how to measure whether the intervention is working. It does not set final list prices or approve individual account tactics. Those decisions require competitive context, contract terms, renewal dates, and account intelligence that sit outside the transaction data.

Decision need	Evidence used in this report	Management output
Size the value pool	List value, billed revenue, realization trend, annual leakage	Realization-improvement target
Locate the exposure	Segment, channel, category, region and segment-channel cuts	Priority plays by cell
Turn diagnosis into action	Customer risk scores, tiering, concentration and driver analysis	Named account worklist

Figure 1

One in five list-price dollars is given away every month

Share of list value actually billed. The shaded band is revenue forgone to discount



Source: pricing governance pipeline, processed marts. January 2023 to December 2025.

Figure 1. Price realization has held inside a two-point band for three years. The shaded area is the share of list value forgone to discount each month, the standing cost this review quantifies.

3 Data and methodology

Every figure in this report traces to a single reproducible pipeline that moves raw transactions through cleaning, feature construction, a SQL warehouse layer, risk scoring, and validation. The pipeline reads five raw tables, customers, products, orders, order items, and sales representatives, and resolves them into an enriched order-item grain of 38,391 rows. That grain is the atom of the analysis: one row per product sold on one order, carrying its list price, its realized price, the discount between them, an estimated unit cost, and the segment, channel, region, and tenure context of the sale.

Layer	Role in the analysis	Primary control
Raw inputs	Customers, products, orders, order items and sales representatives	Schema and required-column checks
Order-item grain	Joins commercial context to realized price, list price, cost and discount	Row-count and grain integrity checks
Feature layer	Constructs discount depth, margin proxy, residual price dispersion and exposure flags	Metric-contract reconciliation
Scoring layer	Prioritises customers by dependency, margin erosion and pricing risk	Tier and score distribution checks
Report layer	Publishes only validated metrics and rendered chart assets	Release gate before PDF generation

The margin proxy and why it is a proxy

Cost of goods is estimated rather than booked, so margin throughout this report is a proxy: realized revenue less estimated cost, divided by realized revenue. That boundary matters. The proxy is strong for like-for-like comparison because every segment, channel and category is measured on the same basis. It is not a statement of accounting gross margin and is not used as one. The recommendations depend on relative gaps, such as Enterprise versus SMB or Professional Services versus the portfolio, not on the absolute margin level.

Key metric definitions

Four constructed metrics carry most of the analytical weight, and each is defined once here so the findings can be read without ambiguity.

Metric	Definition
Realized discount	1 minus realized price over list price, per line; aggregated revenue-weighted unless stated.
Price realization	Realized revenue over list-price revenue. The complement of the weighted discount.
Margin proxy	Realized revenue less estimated cost, over realized revenue.
High-discount line	A line is high-discount at or above a 20% threshold; sensitivity to this cut is tested in Section 9.
Governance priority	A 0-100 composite of discount dependency, margin erosion, and pricing-risk scores per customer.

The customer-level scores deserve a word because the priority queue depends on them. Each customer receives a discount-dependency score, a margin-erosion score, and a pricing-risk score, each on a 0-100 scale, blended into a single governance-priority score and reliability-weighted by how much data the customer provides. A customer with six orders and a customer with seventy are not scored as if equally certain. The blended score then sorts customers into four tiers with a recommended action each, which is the operational output the commercial team works from.

Pipeline integrity and validation

The pipeline is deterministic and seeded, so a fixed environment reproduces the same outputs on repeated runs, and a release gate runs a battery of validation checks before any figure is published. For this review, all ten formal-analysis validation checks passed, and the metric-contract layer confirmed that every published metric matched its independently recomputed value within tolerance. The data is synthetic, generated to be internally consistent rather than drawn from a live system. Section 10 treats that as the main boundary condition. The report should therefore be read as a decision-support demonstration of a repeatable governance method, not as a claim about a real company's historical performance.

4

Analytical framework

The single question, healthy growth or discount-reliant growth, is answered along four axes, and the findings sections follow them in order. Each axis isolates one way that a discounting problem can hide.

The aggregate axis asks whether the book as a whole realizes its list prices. It is answered by price realization and the revenue forgone to discount, read over time so that a stable problem can be told apart from a deteriorating or improving one. A single snapshot cannot make that distinction; 36 months can.

The depth axis asks what discounting costs at the margin. It is answered by the relationship between discount depth and margin, measured two ways as a cross-check: across discount buckets at the line grain, and across customers at the account grain. Agreement between the two supports a consistent descriptive reading of how discount depth and margin move together.

The structural axis asks where the discounting lives. It is answered by cutting discount and margin by segment, channel, region, and product category, and by the interaction of segment and channel, because a problem that is uniform requires a blunt policy response while a problem that is concentrated requires a targeted one. The data, as Section 7 shows, is firmly the second case.

The account axis asks who is responsible for the exposure. It is answered by the customer risk scores, the priority tiers, and the concentration of revenue, which together turn a diagnosis into a worklist. This is the axis that makes the review operational rather than only descriptive.

A fifth consideration, sensitivity, runs underneath all four. Every threshold in the analysis, above all the 20% high-discount line, is tested by moving it and watching whether the conclusion moves with it. Section 9 shows it does not across 15%, 20%, 25%, which is what distinguishes a finding from an artefact of where a cutoff happened to fall.

5

Findings: the aggregate picture

Start with the whole book. Over January 2023 to December 2025 the business billed \$1.88B against a list value of \$2.30B. The \$419M difference is the cost of discount, and at 81.8% realization it is large in both absolute and relative terms. For every five dollars the catalogue says the book is worth, it collects a little over four.

What turns this from a number into a finding is its stability. Monthly price realization sits in a 80.9%-to-82.5% band for the entire window, a spread of barely two points across three years that include whatever seasonality, mix shift, and deal cycles the period contained. The weighted discount behind it moves just as little, oscillating around 18.2% without trend. Revenue, plotted against that flat discount line in Figure 2, is similarly trendless. The book is not growing into its discount or out of it. It is holding both constant.


Year	Revenue	List value	Forgone	Realization	High-disc rev.
2023	\$598M	\$731M	\$133M	81.8%	33.4%
2024	\$641M	\$784M	\$143M	81.8%	33.2%
2025	\$644M	\$787M	\$143M	81.8%	33.5%

The annual view is included to separate signal from monthly noise. The same pattern holds after seasonality is absorbed into full-year totals: realization remains in the low eighties, forgone revenue remains material, and high-discount revenue remains a recurring part of the book. The issue is therefore not a calendar effect; it is a commercial operating pattern.

Figure 2

Revenue held flat while discount pressure stayed elevated

Monthly billed revenue (bars) against the revenue-weighted realized discount (line), January 2023 to December 2025



Source: pricing governance pipeline, processed marts, January 2023 to December 2025.

Figure 2. Monthly revenue against the revenue-weighted discount. Neither trends. The discount is a level the organisation operates at, not a response to conditions that vary.

The operating interpretation is that a purely tactical discount would normally move with deal cycles, mix, or seasonality. This one barely moves. That does not prove a formal policy is causing the level, but it does show that the business is operating around a stable list-to-net norm. Stable norms can be governed: target them, change the approval path, and track whether the level resets.

It also changes how the opportunity should be sized. The \$419M gap is measured on the analysed period, not assumed recoverable in full. A two-point improvement in realization, from 81.8% toward 83.8%, is worth about \$46M on the analysed list base, or \$15M per run-rate year at current scale. That is the value yardstick for the governance programme.

Realization lift	New realization	Revenue capture on analysed list base
+1 pp	82.8%	\$23M
+2 pp	83.8%	\$46M
+3 pp	84.8%	\$69M

These scenarios are not forecasts and do not assume demand elasticity, competitive response, or renewal timing. They are sizing cases: what management would recover if the existing list-price base billed one, two, or three points closer to list. The cases create a practical threshold for prioritising governance work and deciding how much commercial disruption is worth testing.

6 Findings: what discount costs the margin


The aggregate tells us discount is large and stable. The depth axis tells us what it is associated with. The answer is commercially important: deeper discount and thinner margin appear together whether the data is read by order line or by customer.

At the line grain, sorting all 38,391 order items into discount bands and computing blended margin within each produces a near-monotone staircase. The shallowest band earns a 51.7% margin proxy; the deepest earns 38.0%. The steps do not flatten at the bottom, where management would want evidence that extra depth is buying something valuable enough to offset the concession. That evidence is not visible in the transaction data.

Figure 3

Each step deeper into discount strips roughly four margin points

Blended gross-margin proxy within each discount band



Source: pricing governance pipeline, processed marts. January 2023 to December 2025.


Figure 3. Blended margin proxy falls at every step into deeper discount. The relationship is monotone, with no flattening at depth to suggest the deepest discounts buy compensating value.

The same gradient appears one level up, at the customer grain, and this is what makes the reading useful. Plotting each of the 1,173 customers by average discount against average margin yields a correlation of -0.68, with margin falling by 0.96 points for every additional point of discount. That is a warning signal: at account grain, deeper-discounted customers are also materially thinner-margin customers. Any argument that depth is being offset by volume, mix or strategic value needs evidence outside the current transaction data.

Figure 4

Deeper discounting reliably tracks thinner margin

Each bubble is one of 1,173 customers, sized by revenue



Source: pricing governance pipeline, processed marts, January 2023 to December 2025.

Figure 4. Every customer plotted by discount and margin, sized by revenue. The downward fit is steep and tight; deeper-discounted customers are reliably thinner-margin customers, with no volume offset visible.


Two cuts of the same data agreeing this closely support a strong management hypothesis, not a controlled causal claim. The mechanical link is direct: discount is subtracted from price, while cost does not fall because the deal was discounted. Unless a deal brings enough compensating value through volume, retention, or strategic account access, depth comes out of margin. The transaction data does not show that offset at scale.

Composition confirms the depth is not confined to a harmless tail. Bucketing revenue rather than counting lines, the bulk of the book sits in the 10-to-20% band, but the bands above the high-discount line carry real weight, and at the 20% cut 33.4% of all revenue flows through deals discounted past it. A third of the book is in the part of the discount distribution where, per Figure 3, margin is measurably impaired.

Figure 5

Most revenue carries a 10-20% discount, but the deep tail is heavy

Billed revenue by discount band. Bands above 20% are highlighted




Source: pricing governance pipeline, processed marts. January 2023 to December 2025.

Figure 5. Revenue by discount band. The 10-20% band dominates, but the highlighted deep bands are not a rounding error; they carry roughly a third of revenue.

Figure 6

Discounting clusters in a tight 10-20% band, not at the edges

All 38,391 order items by realized discount depth



Source: pricing governance pipeline, processed marts. January 2023 to December 2025.

Figure 6. The distribution of line-level discount. Discounting clusters tightly rather than spreading, another sign that list-to-net practice is operating around common defaults rather than purely case-by-case negotiation.

7 Findings: where the discounting lives

A problem that is everywhere needs a different remedy than a problem that is somewhere. This section cuts the discount four ways to establish which it is. The answer is consistent across every cut: the discounting is concentrated, in the Enterprise segment, in the indirect channels, and in one product category, while it is essentially uniform across geography. Concentration is what makes a targeted response possible.


Segment is the primary axis of concentration

The four segments form a clean ladder. Enterprise discounts deepest at 20.1% and earns the thinnest margin at 43.4%; SMB discounts shallowest at 11.9% and earns the fattest margin at 52.0%, with Mid-Market and Public Sector ranged between. The ladder would be unremarkable if the segments were small. They are not. Enterprise is \$1.10B, 58.6% of the entire book, which means the worst pricing position by margin is also the largest by revenue. That is the structural fact management should act on first.

Figure 7

Enterprise carries the most revenue at the worst discount-margin position

Segment position by discount and margin. Bubble area is total revenue



Source: pricing governance pipeline, processed marts. January 2023 to December 2025.

Figure 7. Segments positioned by discount and margin, sized by revenue. The largest bubble sits in the worst corner: Enterprise combines the most revenue with the deepest discount and the thinnest margin.

Segment	Revenue	Avg discount	High-disc share	Margin proxy
Enterprise	\$1.10B	20.1%	49.6%	43.4%
Public Sector	\$154M	18.0%	27.4%	42.6%
Mid-Market	\$429M	16.0%	11.9%	48.1%
SMB	\$196M	11.9%	1.2%	52.0%

The high-discount share column is where the concentration is starkest. In Enterprise, 49.6% of revenue is deeply discounted; in SMB it is 1.2%. These are not two points on a smooth gradient but two different pricing regimes operating in the same company. SMB provides a lower-risk internal benchmark for the deeper discount pattern observed in Enterprise.

Segment	Forgone	Share of forgone	Realization	Customers
Enterprise	\$279M	66.4%	79.8%	247
Public Sector	\$34M	8.0%	82.0%	60
Mid-Market	\$81M	19.3%	84.1%	344
SMB	\$26M	6.2%	88.2%	522

The value-pool view confirms the operational priority. Enterprise is not only the largest revenue segment; it is also the largest contributor to discount leakage, carrying 66.4% of total forgone revenue. That is why the recommended programme starts with Enterprise account renewals rather than broad enablement or generic sales training. The problem is measurable at segment level and actionable at account level.

Channel sharpens the same picture

Cutting by channel shows the indirect routes to market discount markedly deeper than the direct ones. The reseller channel runs a 21.0% weighted discount and the partner channel 19.3%, against 18.1% Direct and only 14.4% Online. Direct is the largest channel at \$1.06B, so its discount sets the book average, but the indirect channels are where depth concentrates per dollar of revenue.


Channel	Revenue	Forgone	Realization	Margin
Reseller	\$198M	\$53M	79.0%	43.4%
Partner	\$381M	\$91M	80.7%	44.5%
Direct	\$1.06B	\$234M	81.9%	45.1%
Online	\$246M	\$41M	85.6%	48.8%

The channel value pool matters because an approval rule changes behaviour only where it is placed. Direct carries the largest dollars, but Reseller and Partner carry the lowest realization and the highest depth. The control therefore belongs at the indirect-channel deal review point: it catches the behaviour with the worst realization without slowing down every direct renewal.

Figure 8

Reseller discounts run nearly 50% deeper than Online

Revenue-weighted discount by channel. Revenue in white at base



Source: pricing governance pipeline, processed marts. January 2023 to December 2025.


Figure 8. Weighted discount by channel, revenue noted at each base. Indirect channels discount close to half again as deep as Online.

The interaction of segment and channel is where the exposure peaks, and it is worth isolating because it names the single worst cell in the book. Where Enterprise meets the reseller channel, average discount reaches 23.4% and 87.9% of that cell's revenue is deeply discounted, against a margin proxy of 40.8%. The heatmap makes the gradient legible: discount darkens reliably from the SMB-Online corner to the Enterprise-reseller corner, the two regimes sitting at opposite ends of the same grid.

Figure 9

Discounting concentrates where Enterprise meets the reseller channel

Average discount by segment and sales channel. Darker is deeper



Source: pricing governance pipeline, processed marts. January 2023 to December 2025.

Figure 9. Average discount by segment and channel. The gradient runs cleanly from the disciplined SMB-Online corner to the Enterprise-reseller corner, which is the deepest cell in the book.


Geography is not the primary discriminator

Region is the control cut. Weighted discount by region produces little spread: APAC is deepest at 18.8% and North America is shallowest at 18.0%, a 0.8% range across four regions that differ materially in size. Geography is therefore not the primary explanation. Discount depth travels more clearly with who is being sold to and through which channel.

Figure 10

Discount depth is roughly even across regions

Weighted discount and revenue by customer region



Source: pricing governance pipeline, processed marts. January 2023 to December 2025.

Figure 10. Weighted discount by region. The narrow spread makes region a weaker discriminator than segment and channel in this dataset.


One product category breaks the pattern

Product category mostly tracks margin as expected, with one exception sharp enough to act on. Professional Services discounts at 17.2%, in line with the rest of the catalogue, but earns a margin proxy of only 20.6%, less than a third of the 62.9% that Collaboration earns and far below the 45.3% book average. It is discounted like a high-margin product while behaving like a low-margin one. Core Platform, by contrast, is the revenue engine at \$925M and earns a respectable 45.3%, so the category problem is specific rather than general.

Figure 11

Professional Services earns the thinnest margin on an ordinary discount

Discount versus margin proxy by product category



Source: pricing governance pipeline, processed marts. January 2023 to December 2025.

Figure 11. Discount against margin by category. Professional Services is the outlier: ordinary discount, exceptionally thin margin, which makes its discount the least defensible in the book.


Concentration cuts the right way

The structural findings would be discouraging if the exposure were spread thinly across thousands of accounts. It is not. Revenue is heavily concentrated: the top 10% of customers produce 46.3% of revenue and the top 20% produce 66.3%. The same concentration that makes the book sensitive to a few accounts also makes it tractable. A governance programme that engages the largest few hundred relationships reaches the great majority of both the revenue and the discount, which is the bridge from this section's diagnosis to the account-level worklist in the next.

Figure 12

Revenue is highly concentrated, which makes targeted remediation viable

Lorenz curve of revenue across the customer base




Source: pricing governance pipeline, processed marts. January 2023 to December 2025.

Figure 12. The Lorenz curve of revenue. Steep early rise means high concentration, and high concentration means targeted remediation can reach most of the exposure.

Figure 13

The top-revenue product also carries the deepest discount exposure

Top 12 products by revenue. Red marks any product with 30%+ of revenue deeply discounted



Source: pricing governance pipeline, processed marts. January 2023 to December 2025.

Figure 13. Top products by revenue, flagged where deep-discount exposure is high. Large products with elevated deep-discount share deserve explicit product-level guardrails.

8

Findings: who carries the risk

The final axis turns the structural picture into a list of names. The risk-scoring layer assigns every customer a governance-priority score and sorts the book into four tiers, each with a recommended action. This is the output the commercial organisation can act on directly.

The tiers are deliberately unbalanced because the risk is. The critical and high tiers hold just 66 customers but \$267M of revenue, at an average priority score of 71 out of 100, all flagged for discount-term review. The medium tier is 168 customers and \$624M, marked for segment-pricing review. The low tier is the long, healthy tail: 939 customers and \$992M that need only monitoring. The shape of the priority distribution, a thin upper tail above a broad healthy base, is what makes a small high-touch programme viable.


Tier	Customers	Revenue share	Avg discount	Avg margin	High-disc rev.
Critical	2	0.5%	24.2%	40.1%	99.0%
High	64	13.7%	22.2%	41.8%	76.6%
Medium	168	33.1%	19.5%	43.4%	44.9%
Low	939	52.7%	13.9%	50.2%	5.7%

The tier table is the bridge from analytics to account management. The critical and high tiers are smaller, more discounted, and more exposed to high-discount revenue than the rest of the book, so it should be handled as a named-account intervention. The medium tier is too large for bespoke treatment; it needs rule-based review as renewals come up. The low tier should not be disturbed unless its score moves. This prevents governance from becoming a blanket control that slows healthy business.

Figure 14

A small high-risk group sits on outsized revenue

Customers and revenue by governance risk tier




Source: pricing governance pipeline, processed marts. January 2023 to December 2025.

Figure 14. The four risk tiers by customer count and by revenue. A few dozen high-risk accounts carry revenue out of all proportion to their number.

Figure 15

Priority is concentrated in a thin upper tail

Governance priority score across 1,173 customers



Source: pricing governance pipeline, processed marts. January 2023 to December 2025.

Figure 15. The governance-priority distribution. Most customers score low; the population that needs action sits above the 90th percentile at a score of 57.


What is actually driving the scores

The scoring layer also records why each customer is flagged, which prevents the queue from being a black box. Aggregating the primary driver across the book, discount dependency is dominant: it is the lead driver for 455 customers carrying \$1.38B, far more revenue than the pricing-risk and margin-erosion drivers behind it. This matters for the remedy. A book whose main problem is dependency, customers who have come to expect a standing discount, responds to renegotiation and approval discipline. It would not respond to the same degree if the driver were, say, cost inflation, which no amount of pricing governance would touch.

Figure 16

Discount dependency is the dominant driver by revenue

Primary risk driver assigned to each customer, aggregated by revenue



Source: pricing governance pipeline, processed marts, January 2023 to December 2025.

Figure 16. Revenue by primary risk driver. Discount dependency dominates, which points the remedy squarely at discount policy rather than cost or mix.


Discount depth persists across tenure cohorts

One more cut closes the loop between this axis and the aggregate finding. Splitting customers by tenure shows discount depth essentially flat across the lifecycle: newly acquired customers and four-year veterans are discounted at materially the same rate. The cross-sectional pattern is consistent with persistent discounting, but it does not track the same account longitudinally and therefore cannot distinguish acquisition incentives from contract, mix, or cohort effects on its own.

Figure 17

Discount depth is similar across customer-tenure cohorts

Average discount and margin proxy by customer tenure band



Source: pricing governance pipeline, processed marts. January 2023 to December 2025.

Figure 17. Discount and margin by customer tenure. Similar discount levels across tenure bands support a persistence hypothesis that should be tested with longitudinal renewal data.

9 Sensitivity and confidence


The most consequential assumption in the review is the high-discount line. It determines high-discount revenue share and feeds the risk scores. A credible conclusion should not disappear when that line moves, so this section tests the verdict across 15%, 20%, 25%.

At a 15% line, 76.9% of revenue counts as high-discount and \$162M of revenue is flagged as margin-at-risk. At the 20% line used throughout, those fall to 33.4% and \$65M. At 25%, to 4.3% and \$13M. The magnitudes move, as they must, but the direction and the diagnosis do not. At every threshold the same segments, channels, and accounts surface as the deepest, and at every threshold the margin proxy on high-discount revenue sits below the book average. The cutoff changes the size of the worklist; it does not change where management should start.

Figure 18

The risk verdict remains stable across the tested thresholds

Revenue at risk (bars) and high-discount share (line) as the threshold moves



Source: pricing governance pipeline, processed marts. January 2023 to December 2025.

Figure 18. Revenue at risk and high-discount share as the threshold moves from 15% to 25%. The curve is smooth and the ranking of risk is stable, so no single cutoff manufactures the result.

Threshold	High-disc revenue share	Revenue margin-at-risk	Margin on high-disc revenue
15%	76.9%	\$162M	44.1%
20%	33.4%	\$65M	42.9%
25%	4.3%	\$13M	41.1%

Two separate validation layers add to the confidence. The metric-contract layer recomputes every published figure from the raw grain and checks it against the value the pipeline reports, and all checks reconciled within tolerance. The formal-analysis release gate ran ten structural checks on the analysis window, coverage, and internal consistency, and all ten passed. The numbers in this report are therefore recomputed and reconciled before publication rather than copied from an output table.

Challenge to the finding	Check performed	Result
Threshold dependence	Moved high-discount line across 15%, 20%, 25%	Diagnosis and ranking remain stable
Aggregation artefact	Compared line-level bucket results with customer-level scatter	Both show the same discount-margin gradient
Geographic mix	Cut weighted discount by region	Regional spread is under one point
Single-driver risk	Aggregated primary score drivers by customer revenue	Discount dependency is the dominant driver
Metric publication error	Recomputed published metrics through contracts and release gate	All checks passed

The confidence case is therefore cumulative rather than dependent on any one statistic. The same commercial story appears in the aggregate trend, in the line-level margin staircase, in the customer scatter, in the segment-channel heatmap, and in the customer risk queue. A different threshold would resize the opportunity, but it would not change where management should start.

10 Risks, limitations, and caveats

The analysis is useful because it is bounded. Four limitations matter, and each changes how the work should move from diagnostic to execution.

Boundary	What it could change	Evidence needed to close
Synthetic data	Absolute dollar opportunity and account names in a live setting	Production order, cost and contract data
Margin proxy	Level of profitability by product or account	Booked cost, implementation effort and service delivery cost
Observational design	Causal size of a price-realization intervention	Renewal test, matched cohort or controlled price experiment
Score weighting	Borderline medium-versus-low prioritisation	Commercial review of accounts near tier thresholds

The data is synthetic

The order book is generated rather than drawn from a production system. It is internally consistent and behaves like a real book, which makes it suitable for validating the method, but the dollar figures are not a claim about any real company. The value of the review is the operating model: metrics, scoring, validation, and a decision structure that can be pointed at a live order book and re-run with production data.

Margin is a proxy

Because cost is estimated, margin is a proxy and is used only for comparison, never as a statement of absolute profitability. The comparisons are internally consistent because every group uses the same cost-generation method, but actual booked cost could change both levels and rankings. The margin findings must therefore be revalidated before operational use with production economics.

Correlation is not full causation

The near-one slope between discount and margin is consistent with discount driving margin loss, and the agreement between the line and customer cuts strengthens that reading. It is still observational evidence. A controlled renewal test or matched cohort is needed to quantify how much realization can be recovered without affecting win rate, renewal risk, or customer expansion.

The scores encode judgement

The governance-priority score blends three sub-scores with chosen weights, and different weights would reshuffle the queue at the margin. The tiering is therefore a defensible prioritisation rather than an objective truth, and it should be used as a starting worklist that commercial judgement refines, not as an automated verdict on any single account. Weight sensitivity should be tested during calibration, especially for accounts close to a tier boundary.

11 Recommendations and action priorities

The response should be run as a 90-day price-realization recovery programme, not as a generic pricing initiative. The objective is narrow: test whether management can move realization from 81.8% toward 83.8% in the exposed parts of the book while protecting renewal and win-rate. That two-point move is worth about \$46M on the analysed list base, or \$15M per year at current scale.

Priority 1: ring-fence the high-risk accounts for renewal renegotiation

The 66 critical and high-risk accounts carry \$267M at an average priority of 71 and are individually nameable from the scoring output. Assign each account to an executive owner and a pricing owner, set a target net price before negotiation starts, and require explicit approval for any renewal that lands below target. Because dependency is the dominant driver, the realistic objective is to narrow the discount by a few points, not to remove it. The account score should become a renewal workplan: target price, concession logic, exception owner, and realized outcome.

Priority 2: put an approval gate on deep discounts in the indirect channels

Depth concentrates in the reseller and partner channels, peaking in the Enterprise-reseller cell at 23.4% with nearly nine-tenths of that revenue deeply discounted. Introduce a hard approval gate above the 20% line for these channels. Every exception should carry a reason code, expected account value, margin check, and approval owner. This is the forward-looking control that stops the standing level from regenerating after Priority 1 resets individual renewals.

Priority 3: re-price Professional Services

Professional Services is discounted like the rest of the catalogue, at 17.2%, while earning a 20.6% margin proxy against a book average above 45.3%. This is the least defensible discount in the portfolio. Separate services authority from product discount authority, set a services-specific floor, and stop treating labour-heavy delivery as if it had software margin. If the market will not bear the required net price, the business should decide that explicitly as a margin trade, not hide it inside a product discount bundle.

Priority 4: adopt price realization as a standing operating metric

The reason the \$419M accumulated unnoticed is that no one owned realization as a number. Put it beside revenue in the monthly business review, report it by segment and channel, and hold commercial owners accountable for both volume and quality of revenue. The measurement is already built; the pipeline produces realization at every cut in this report. What is missing is ownership.

Control model

Control point	Rule	Review cadence	Owner
High-risk renewals	Target net price set before negotiation; variance from target explicitly approved	Renewal calendar	Segment GM + Finance
Indirect deep discounts	Deals above 20% require reason code and margin review	Deal desk	Channel Sales + Pricing
Professional Services	Separate service pricing guardrail from product discount authority	Quarterly	Services GM + FP&A
Realization monitoring	Monthly realization by segment, channel and category; escalation on adverse movement	Monthly business review	Revenue Operations

The control model is intentionally light. It does not ask management to police every order line. It places scrutiny where the analysis shows repeatable exposure: high-risk renewals, indirect-channel deep discounts, and the low-margin services category. The monthly realization metric then tests whether the controls are changing the book, not just individual approval behaviour.

90-day mobilisation plan

Window	Management actions	Exit criterion
Days 0-15	Freeze the 66-account list, assign owners, set renewal targets, and publish the 20% indirect-channel approval rule	Named worklist and approval path live
Days 16-45	Run the first renewal wave, collect exception reason codes, and design the Professional Services floor	First-wave realization lift and exception mix reviewed
Days 46-90	Expand to the medium tier as renewals arrive, launch services guardrail, and embed realization in the monthly business review	Book-level realization tracking toward 83.8%

Further questions before scale

- **Elasticity and win-rate:** which Enterprise and indirect-channel deals can absorb a two-to-three point net-price lift without materially reducing close rate?
- **Contract timing:** which of the high-risk accounts renew inside the next two quarters, and what share of the value pool is addressable in the current planning cycle?
- **Competitive exceptions:** which deep discounts are tied to documented competitive displacement, regulated procurement rules, or strategic account-entry logic?
- **Cost calibration:** how do booked cost and implementation effort change the Professional Services margin proxy once actual delivery economics are loaded?

These questions do not weaken the recommendation; they define the next evidence needed to turn a diagnostic programme into a priced execution plan. The current analysis is sufficient to prioritise where to act. The open questions determine how much of the value pool can be captured, when, and with what commercial trade-off.

12 Appendix

A. Segment detail

Segment	Revenue	Avg discount	High-disc share	Margin proxy
Enterprise	\$1.10B	20.1%	49.6%	43.4%
Public Sector	\$154M	18.0%	27.4%	42.6%
Mid-Market	\$429M	16.0%	11.9%	48.1%
SMB	\$196M	11.9%	1.2%	52.0%

B. Channel detail

Channel	Weighted discount	Revenue
Reseller	21.0%	\$198M
Partner	19.3%	\$381M
Direct	18.1%	\$1.06B
Online	14.4%	\$246M

C. Product category detail

Category	Revenue	Discount	Margin proxy
Core Platform	\$925M	18.4%	45.3%
Security	\$364M	18.8%	47.8%
Analytics	\$258M	18.4%	52.6%
Professional Services	\$205M	17.2%	20.6%
Collaboration	\$130M	16.6%	62.9%

D. Region detail

Region	Revenue	Weighted discount	Margin proxy
North America	\$812M	18.0%	45.4%
Europe	\$499M	18.1%	45.6%
APAC	\$322M	18.8%	44.7%
LATAM	\$250M	18.4%	45.3%

E. Risk tier detail

Tier	Recommended action	Customers	Revenue	Avg priority
Critical	review discount terms	2	\$9M	81
High	review discount terms	64	\$258M	71
Medium	review segment pricing	168	\$624M	53
Low	monitor only	939	\$992M	27

Provenance. All figures derive from the processed marts of the pricing-governance pipeline over January 2023 to December 2025: 38,391 order lines, 1,173 customers, 28 products. Ten formal-analysis validation checks and the full metric-contract reconciliation passed at the time of publication. Charts are generated by scripts/build_report_assets.py and this report by scripts/build_report_pdf.py, both reproducible from the processed data.